

UNIVERSITY OF POTSDAM

BACHELOR THESIS

Investigation of Functionals of the Eigenvalues of Unitary Matrices

Author:
Carina SEIDEL

Supervisor:
Dr. Thomas MACH

July 20, 2025



Abstract

The eigenvalues of large matrices are of interest in a large variety of use cases. However, their computation becomes increasingly challenging as matrix size grows. This thesis focuses on the investigation of functionals over the eigenvalues of unitary matrices, with particular emphasis on the spectral density, or density of states (DoS) among those. While established methods exist for real symmetric matrices, this work extends these approaches to unitary matrices with the help of the Cayley transform. The aim is to develop and analyze efficient techniques for approximating spectral functionals in this broader context.

Contents

1	Introduction	2
1.1	Unitary Matrices	2
1.2	Cayley Transform	4
1.3	Spectral Density	5
2	Motivation	7
3	Kernel Polynomial Method	8
3.1	Overview	8
3.2	Polynomial Extension with Chebyshev polynomials	9
3.3	Approximating the spectral density	10
3.4	The Choice of Random Unitary Matrices	12
4	Quality Analysis	14
4.1	Resolution	14
4.2	Restriction on the Schwartz space	14
4.3	Regularizing the Dirac delta distribution	15
4.4	The condition of non-negativity	17

Introduction

The importance of eigenvalues in mathematics and physics is well established. However, for very large matrices, computing all eigenvalues can be prohibitively expensive. In many applications, it is sufficient to know only a few extremal eigenvalues, or even just the overall distribution of eigenvalues. This leads to the concept of the spectral density, which describes how eigenvalues are distributed for a given matrix.

While the spectral density is well understood for real symmetric matrices, our goal is to extend this concept to the complex case. To do so, we first examine Hermitian matrices, which generalize symmetric matrices to the complex domain. We then introduce the Cayley transform, which serves as a bridge to unitary matrices. Finally, we define and investigate the spectral density for unitary matrices, which plays a central role in many areas of mathematics and physics.

We begin this thesis by recalling fundamental properties of unitary matrices. Next, we introduce the Cayley transform to connect Hermitian and unitary matrices, and then approach the concept of spectral density. With this foundation, we are well-equipped to proceed with the main investigation.

1.1 Unitary Matrices

We first recall two definitions for important real matrices that we then extend to complex matrices. The index T marks the transpose of a matrix. As common in literature, I_n denotes the identity matrix of size n .

Definition 1 (Symmetric matrix). Let A be a real, square matrix of size n . Then A is called *symmetric* if $A^T = A$.

The following definition is for matrices with their transpose as their inverse.

Definition 2 (Orthogonal matrix). Let A be a real, square matrix of size n . Then A is called *orthogonal* if $A^T \cdot A = A \cdot A^T = I_n$.

Now we define the complex equivalent of a real symmetric matrix. Here, A^* denotes the *conjugate transpose* (also called adjoint) of A , that is, the matrix obtained by taking the transpose and then complex conjugating all entries.

Definition 3 (Hermitian matrix). Let A be a complex square matrix of size n . Then A is called *Hermitian* if $A^* = A$.

Throughout this thesis, A will denote a complex, square matrix of size n , unless

stated otherwise. When referring to a Hermitian matrix, we will use the letter H .

We now examine the eigenvalues of Hermitian matrices.

Let $H = H^*$ and $Hv = \lambda v$ for a complex vector $v \neq \mathbf{0}$ of size n and a scalar $\lambda \in \mathbb{C}$. Consider now the inner product v^*v . Remember that $(A \cdot B)^* = B^* \cdot A^*$ holds for a matrix product, and obviously $(A^*)^* = A$ for all matrices and vectors.

$$\lambda v^*v = v^*(\lambda v) = v^*(Hv) = (v^*H)v = (H^*v)^*v = (Hv)^*v = (\lambda v)^*v = \bar{\lambda}v^*v$$

Since we have that $v \neq \mathbf{0}$ it follows that $v^*v \neq \mathbf{0}$ and therefore $\lambda = \bar{\lambda}$, that is to say λ is real. This means that all eigenvalues of Hermitian matrices are real numbers. It follows that all eigenvalues of symmetric matrices are also real numbers, since they are a special case of Hermitian matrices. Now we can define the complex equivalent of orthogonal matrices:

Definition 4 (Unitary matrix). A matrix A is called *unitary* if $A^* \cdot A = A \cdot A^* = I_n$.

We will oftentimes denote unitary matrices by using U as a reference. It is easy to see that orthogonal matrices are a special case of unitary matrices, since $A^T = A^*$ for all real matrices.

Consider a unitary matrix U and an eigenpair (λ, v) of U . The complex conjugate of the eigenvalue equation $Uv = \lambda v$ is

$$v^*U^* = v^*\bar{\lambda} = \bar{\lambda}v^*$$

We calculate

$$v^*v = v^*I_nv = v^*U^*Uv = v^*\bar{\lambda}\lambda v = \bar{\lambda}\lambda v^*v = |\lambda|^2 v^*v$$

Similarly to above, we can divide by v^*v to obtain

$$1 = |\lambda|^2 \implies |\lambda| = 1$$

meaning that all eigenvalues of unitary matrices have a length of 1 and are thus situated on the unit circle. As they are a special case of unitary matrices, the same goes for orthogonal matrices. This property is crucial, as it enables the application of the Cayley transform, introduced in the following section.

There is a special group of matrices that all the matrices we have defined so far belong to.

Definition 5 (Normal matrix). A matrix A is called *normal* if it commutes with its conjugate transpose, that is to say $A^*A = AA^*$.

It is straightforward to see that both Hermitian and unitary matrices are normal matrices and that the notion includes real symmetric and orthogonal matrices as special cases.

The spectral theorem states that normal matrices can be diagonalized by a unitary matrix. That means, for any normal matrix A , there exists a unitary matrix U such that $A = U\Lambda U^*$, where $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$ with $\lambda_1, \dots, \lambda_n$ being the eigenvalues of A . This is essential for applying a function to a matrix, which we will get to in the next section.

1.2 Cayley Transform

Before giving the central definition of this section, we will first extend the concept of functions to normal matrices, allowing us to map one matrix to another.

Definition 6 (Matrix function on normal matrices). Let A be a normal matrix and let $f : \mathbb{C} \rightarrow \mathbb{C}$ be a function that is defined on the spectrum of A , $\sigma(A) = \{\lambda_1, \dots, \lambda_n\}$. Then the *matrix function* $f(A)$ is defined as

$$f(A) := U f(\Lambda) U^* = U \operatorname{diag}(f(\lambda_1), \dots, f(\lambda_n)) U^*$$


where U is the matrix of eigenvectors of A and $\Lambda = \operatorname{diag}(\lambda_1, \dots, \lambda_n)$ is the diagonal matrix of eigenvalues.

Now we can define the *Cayley transform*, which is a specific matrix function, that establishes a correspondence between Hermitian and unitary matrices, allowing spectral properties to be translated between these two important classes.

For a complex number $z \in \mathbb{C}$ with $z \neq -i$, the Cayley transform is defined as

$$\varphi(z) = \frac{i - z}{i + z}.$$

This function maps the real line to the unit circle in the complex plane.



For matrices, the Cayley transform maps a Hermitian matrix H (with $i + H$ invertible) to a unitary matrix U via

$$U = (i - H)(i + H)^{-1}.$$

The condition of $i + H$ being invertible is the same as requiring that H does not have $-i$ as an eigenvalue. If (and only if) that were the case, then we had $Hv = -iv$ for some eigenvector $v \neq \mathbf{0}$, and therefore

$$(i + H)v = iv + Hv = iv + (-iv) = 0.$$

Since Hermitian matrices have only real eigenvalues as discussed above, $-i$ can never be an eigenvalue. Conversely, given a unitary matrix U (with $U \neq -I_n$), the inverse Cayley transform yields a Hermitian matrix:

$$H = i(I_n - U)(I_n + U)^{-1}.$$

This will be relevant later, as we can then use φ to transform unitary matrices into symmetric ones and vice versa.

1.3 Spectral Density

To get to the notion of the spectral density, we will first need some more basic definitions to build on. The reader is assumed to be familiar with the concepts of distributions. Let us also recap that for $\Omega \subset \mathbb{C}^n$ open and non-empty, a *test function* is a smooth function with compact support defined on Ω . The space of all test functions on Ω is usually denoted by \mathcal{E} .

We will now look at an important case of a distribution

Definition 7 (Dirac delta distribution). Let $\mathcal{E} = \mathcal{C}^\infty(\Omega)$ with $0 \in \Omega \subset \mathbb{R}^n$. Then

$$\delta : \mathcal{E} \rightarrow \mathbb{R}, \quad f \mapsto f(0) \quad \text{with} \quad \delta(f) = \langle \delta, f \rangle = f(0)$$

this distribution is often mistakenly referred to as a function, although it is not a function in the classical sense.

The Dirac delta is characterized by the following property:

$$\int_{-\infty}^{\infty} f(x) \delta(x - a) \, dx = \int_{-\infty}^{\infty} f(x) \delta(a - x) \, dx = f(a) \implies \int_{-\infty}^{\infty} \delta(x - a) \, dx = 1.$$

This means that the Dirac delta distribution is zero everywhere except at the point a , where it is infinitely high, such that the integral over it equals 1. We now have all the tools we need to define the central concept of this thesis:

Definition 8 (Spectral density). Let H be hermitian and sparse. For $x \in \mathbb{R}$, the *spectral density* is then defined as

$$\phi(x) = \frac{1}{n} \sum_{j=1}^n \delta(x - \lambda_j)$$

where δ is the Dirac delta distribution and λ_j are the eigenvalues of H in non-descending order.

The summation over the n eigenvalues of H creates a spike at each eigenvalue λ_j , resulting in a distribution that is zero everywhere except at the eigenvalues of H . The division by n ensures that the spectral density is normalized, meaning that the integral over the entire real line equals 1. This is an important property, as it allows us to interpret the spectral density as a probability density function.

The number of eigenvalues ν in an interval $[a, b]$ can then be counted in the following manner:

$$\nu_{[a,b]} = \int_a^b \sum_j \delta(t - \lambda_j) \, dt \equiv \int_a^b n \phi(t) \, dt \quad (1.1)$$

Random matrix theory tells us that the eigenvalues of random unitary matrices are distributed uniformly on the unit circle. Thus, we would expect the spectral density to be uniformly distributed on the unit circle as well. This is important as we think of the choice of random unitary matrices.

If we generate a random complex square matrix A , there are multiple ways to obtain a unitary matrix U . We will now compare the svd with the qr decomposition.

Before we can proceed to the motice of this thesis, we will need to define the following space:

Definition 9 (Schwartz space over \mathbb{R}). The *Schwartz space* over \mathbb{R} consists of all smooth functions f that decay rapidly to zero as $|x|$ approaches infinity [4]. Formally,

$$\mathcal{S}(\mathbb{R}) := \left\{ f \in \mathcal{C}^\infty(\mathbb{R}) \mid \forall p, k \in \mathbb{N}_0 : \sup_{x \in \mathbb{R}} \left| x^p f^{(k)}(x) \right| < \infty \right\}$$

Motivation

Because of its ability to capture the distribution of eigenvalues, the spectral density of a matrix is a crucial concept in various fields, including physics, statistics, and machine learning, although this list is not exhaustive. However, two challenges arise when defining the spectral density of a matrix H :

Firstly, the eigenvalues of the matrix H are typically not known in advance. This would make the computation of the spectral density straightforward, but also defeat the purpose.

Secondly, the delta distribution $\phi(x)$ is not a conventional *function* that can be evaluated pointwise. Therefore, there is a need for efficient methods to approximate the spectral density at low computational cost.

A fundamental idea to address this is to replace the delta distribution by a regular function f with similar properties:

$$\int_{\mathbb{R}} f(x) dx = 1,$$

where $f(x)$ is very small or zero for most $x \neq 0$, and finite but large near $x = 0$. This regularization makes it possible to approximate the spectral density numerically and is central to many practical algorithms.

One intuitive approach is to select an interval $I \subset \mathbb{R}$ such that the spectrum of H , $\sigma(H)$, is contained within I . Next, choose k points t_i in I such that the interval is divided into subintervals:

$$\{t_i\}_{i=1}^k \subset I \quad \text{with} \quad \bigcup_{i=1}^{k-1} [t_i, t_{i+1}] = I$$

Count the number of eigenvalues in each subinterval. Then, calculate the average value of $\phi(x)$ in each interval using $\nu_{[a,b]}$ from equation 1.1. The result is a histogram, which, as the subintervals become smaller (i.e., as k increases and $(t_{i+1} - t_i) \rightarrow 0$), approaches the true spectral density.

To count the eigenvalues in the intervals, one can use methods such as Sylvester's law of inertia. The details of this method are beyond the scope of this work; it would require computing a decomposition of $A - t_i I = LDL^T$ for all t_i [3]. While this can be computationally expensive for general matrices, it remains a viable and efficient option for certain structured matrices. However, for large and unstructured matrices, we prefer methods based on matrix-vector multiplications, which are more efficient in higher dimensions.

Kernel Polynomial Method

3.1 Overview

There exists a whole class of methods, all of which called the kernel polynomial method, or KPM for short. They are powerful tools for approximating the spectral density of matrices. We will focus on the main approach in the following.

As the name suggests, the KPM is a polynomial extension of the spectral density. The coefficients of the polynomials are derived from the method of moments, in order to obtain an estimator function as in statistics.

The method is based on a corollary of the following generalized theorem:

Theorem 1. Let $A \in \mathbb{C}^{n \times n}$ be a normal matrix with spectral decomposition

$$A = U\Lambda U^* \quad \text{where} \quad UU^* = I_n \text{ and } \Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$$

Let $\beta, v \in \mathbb{C}^n$ with $v = U\beta$. Suppose v is a random vector whose entries v_i are independent and identically distributed standard complex normal variables, i.e., $v_i \sim_{\text{i.i.d.}} \mathcal{N}_{\mathbb{C}}(0, 1)$, meaning $\text{Re}(v_i), \text{Im}(v_i)$ are independent $\mathcal{N}(0, \frac{1}{2})$. Then

$$\mathbb{E}[v] = 0 \quad \text{and} \quad \mathbb{E}[vv^*] = I_n, \tag{3.1}$$

and it follows that

$$\mathbb{E}[\beta] = 0 \quad \text{and} \quad \mathbb{E}[\beta\beta^*] = I_n.$$

Proof of Theorem 1. Since the expectation operator is linear, it holds that

$$\mathbb{E}[v] = \mathbb{E}[U\beta] = U\mathbb{E}[\beta] = 0 \implies \mathbb{E}[\beta] = 0$$

Furthermore it holds that

$$I_n = \mathbb{E}[vv^*] = \mathbb{E}[(U\beta)(U\beta)^*] = \mathbb{E}[U\beta\beta^*U^*] = U\mathbb{E}[\beta\beta^*]U^*$$

Multiplying both sides with U^* and U yields:

$$U^*I_nU = U^*U\mathbb{E}[\beta\beta^*]U^*U = \mathbb{E}[\beta\beta^*]$$

Since U is unitary, we have shown that $\mathbb{E}[\beta\beta^*] = I_n$. □

This theorem has a nice corollary when investigating a matrix function $f(A)$. In that case, we have

$$\begin{aligned}
\mathbb{E}[v^* f(A) v] &= \mathbb{E}[(U\beta)^* f(U\Lambda U^*)(U\beta)] \\
&= \mathbb{E}[\beta^* U^* U f(\Lambda) U^* U \beta] \\
&= \mathbb{E}[\beta^* f(\Lambda) \beta] \\
&= \mathbb{E}\left[\sum_{j=1}^n |\beta_j|^2 f(\lambda_j)\right] \\
&= \sum_{j=1}^n f(\lambda_j) \mathbb{E}[|\beta_j|^2] \\
&= \sum_{j=1}^n f(\lambda_j)
\end{aligned}$$

or, more concisely,

$$\mathbb{E}[v^* f(A) v] = \text{Tr}(f(A)). \quad (3.2)$$

We now have a method to calculate the trace of a matrix function $f(A)$, using only vector multiplications with A .

3.2 Polynomial Extension with Chebyshev polynomials

We assume Chebyshev polynomials to be a good fit for the polynomial extension of the Dirac delta distribution because of the many great properties we will see now. First, let's look how Chebyshev polynomials are defined. Using the trigonometric functions, they can be expressed as follows:

$$T_k(x) = \begin{cases} \cos(k \arccos(x)) & \text{for } k \in [-1, 1] \\ \cosh(k \text{ arcosh}(x)) & \text{for } k > 1 \\ (-1)^k \cosh(k \text{ arcosh}(-x)) & \text{for } k < -1 \end{cases}$$

We will only use the formula $T_k(x) = \cos(k \arccos(x))$. This means to only consider matrices, which have eigenvalues within the interval $[-1, 1]$. In the case that this condition should not be fulfilled, the eigenvalues can be transformed accordingly. For this, let λ_{lb} and λ_{ub} be the lower and upper bound for the eigenvalues of A , respectively. To find these, well-established methods like the Gershgorin circle theorem can be used.

Define

$$c := \frac{\lambda_{lb} + \lambda_{ub}}{2} \quad \text{and} \quad d := \frac{\lambda_{ub} - \lambda_{lb}}{2}$$

Then, the matrix $B = \frac{A - cI_n}{d}$ has eigenvalues in the interval $[-1, 1]$. A visualization of this is linked in the appendix.

Another way to define Chebyshev polynomials is by calculating them using the recursion formula

$$T_{k+1}(x) = 2xT_k(x) - T_{k-1}(x),$$

where the starting conditions are given by $T_0(x) = 1$ and $T_1(x) = x$.

Additionally, note that the result in equation 3.2 says, that

$$\mathbb{E} [v^T T_k(A) v] = \text{Tr} (T_k(A)). \quad (3.3)$$

This result will be central in the following discussion.

Now, let

$$h(x) = \frac{1}{\sqrt{1-x^2}} \quad (3.4)$$

be a weight function. Another property of Chebyshev polynomials is that they are *orthogonal* in terms of the scalar product weighted with h :

$$\langle f, g \rangle = \int_{-1}^1 \frac{1}{\sqrt{1-x^2}} \cdot f(x) \cdot g(x) \, dx.$$

This means that

$$\int_{-1}^1 \frac{1}{\sqrt{1-t^2}} \cdot T_k(t) \cdot T_l(t) \, dt = \begin{cases} 0 & \text{for } k \neq l \\ \pi & \text{for } k = l = 0 \\ \frac{\pi}{2} & \text{for } k = l \neq 0 \end{cases}$$

3.3 Approximating the spectral density

Now multiply the spectral density with the inverse of the weight function 3.4:

$$\hat{\phi}(x) = \sqrt{1-x^2} \phi(x) = \sqrt{1-x^2} \cdot \frac{1}{n} \sum_{j=1}^n \delta(x - \lambda_j)$$

Let $g \in \mathcal{S}(\mathbb{R})$, the Schwartz space defined in definition 9, and $\mu_k \in \mathbb{R}$ coefficients to be determined such that the following equation holds:

$$\int_{-1}^1 \hat{\phi}(t) g(t) \, dt = \int_{-1}^1 \sum_{k=0}^{\infty} \mu_k T_k(t) g(t) \, dt \quad (3.5)$$

If this is true for arbitrary $g \in \mathcal{S}(\mathbb{R})$, we can simplify our equation 3.5 to

$$\hat{\phi}(x) = \sum_{k=0}^{\infty} \mu_k T_k(x) \quad (3.6)$$

Now utilize the orthogonality of the Chebyshev polynomials, to calculate a specific coefficient μ_k :

$$\begin{aligned}
& \sum_{l=0}^{\infty} \mu_l T_l(t) = \hat{\phi}(t) \\
\Rightarrow & \left(\sum_{l=0}^{\infty} \mu_l T_l(t) \right) \cdot T_k(t) = \hat{\phi}(t) \cdot T_k(t) \\
\Rightarrow & \int_{-1}^1 \frac{1}{\sqrt{1-t^2}} \cdot \left(\sum_{l=0}^{\infty} \mu_l T_l(t) \right) \cdot T_k(t) dt = \int_{-1}^1 \frac{1}{\sqrt{1-t^2}} \cdot \hat{\phi}(t) \cdot T_k(t) dt \\
\Rightarrow & \mu_k \cdot \frac{\pi}{2 - \delta_{k0}} = \int_{-1}^1 \frac{1}{\sqrt{1-t^2}} \cdot \sqrt{1-t^2} \cdot \phi(t) \cdot T_k(t) dt \\
\Rightarrow & \mu_k = \frac{2 - \delta_{k0}}{\pi} \cdot \int_{-1}^1 \phi(t) \cdot T_k(t) dt
\end{aligned}$$

By applying the Dirac delta distribution we obtain:

$$\begin{aligned}
\mu_k &= \frac{2 - \delta_{k0}}{\pi} \cdot \int_{-1}^1 \phi(t) \cdot T_k(t) dt = \frac{2 - \delta_{k0}}{\pi} \cdot \int_{-1}^1 \frac{1}{n} \sum_{j=1}^n \delta(t - \lambda_j) \cdot T_k(t) dt \\
&= \frac{2 - \delta_{k0}}{n\pi} \sum_{j=1}^n T_k(\lambda_j) \\
&= \frac{2 - \delta_{k0}}{n\pi} \text{Tr}(T_k(A))
\end{aligned}$$

Now let $n_{\text{vec}} \in \mathbb{R}$ and consider vectors $v_0^{(1)}, v_0^{(2)}, \dots, v_0^{(n_{\text{vec}})}$, which are chosen randomly from the standard normal distribution, that is to say $\mathbb{E}[v_0^{(k)}] = 0$ and $\mathbb{E} \left[v_0^{(k)} \left(v_0^{(k)} \right)^T \right] = I_n$. It follows from equation 3.3 that

$$\zeta_k = \frac{1}{n_{\text{vec}}} \sum_{l=1}^{n_{\text{vec}}} \left(v_0^{(l)} \right)^T T_k(A) v_0^{(l)}$$

is a good estimator for $\text{Tr}(T_k(A))$ and therefore

$$\mu_k \approx \frac{2 - \delta_{k0}}{n\pi} \zeta_k$$

In order to determine the ζ_k , let $v_0 \equiv v_0^{(l)}$. Using the recursion formula for Tschebyshev polynomials, we can calculate

$$T_{k+1}(A)v_0 = 2AT_k(A)v_0 - T_{k-1}(A)v_0$$

For $v_k \equiv T_k(A)v_0$ it also holds that

$$v_{k+1} = 2Av_k - v_{k-1}$$

With this we are fully equipped for the final calculation and the goal of the KPM is reached: Instead of having to multiply matrices with other matrices, it now

suffices to multiply matrices with vectors. Now we can approximate $\phi(x)$ as closely as we like. As aforementioned, it is not always desirable to have an infinitely exact approximation. Since it holds that

$$\lim_{k \rightarrow \infty} \mu_k \rightarrow 0$$

and we are only interested in $T_k(x)$ with $k \leq M$
Therefore we estimate ϕ with

$$\tilde{\phi}_M(x) = \frac{1}{\sqrt{1-x^2}} \sum_{k=0}^M \mu_k T_k(x) \quad (3.7)$$

The following pseudo code is based on [2, p. 10] and summarizes the steps described above. The implementation is done in Python, and linked in the appendix.

Algorithm 1 The Kernel Polynomial Method

Require: $A = A^* \in \mathbb{C}^{n \times n}$ with eigenvalues in the interval $[-1, 1]$

Ensure: Estimated spectral density $\{\tilde{\phi}_M(t_i)\}$

```

for  $k = 0 : M$  do
     $\zeta_k \leftarrow 0$ 
end for
5: for  $l = 1 : n_{\text{vec}}$  do
    Choose a random new vector  $v_0^{(l)}$ ;  $\triangleright v_{0_i}^{(l)} \sim \text{i.i.d. } \mathcal{N}(0, 1)$ 
    for  $k = 0 : M$  do
        Calculate  $\zeta_k \leftarrow \zeta_k + \left(v_0^{(l)}\right)^T v_k^{(l)}$ ;
        if  $k = 0$  then
10:  $v_1^{(l)} \leftarrow A v_0^{(l)}$ 
        else
             $v_{k+1}^{(l)} \leftarrow 2A v_k^{(l)} - v_{k-1}^{(l)}$   $\triangleright$  Three term recursion
        end if
    end for
15: end for
    for  $k = 0 : M$  do
         $\zeta_k \leftarrow \frac{\zeta_k}{n_{\text{vec}}}$ 
         $\mu_k \leftarrow \frac{2 - \delta_{k0}}{n\pi} \zeta_k$ 
    end for
20: Evaluate  $\tilde{\phi}_M(t_i)$  with equation 3.7

```

Applying the Cayley transform to an arbitrary unitary matrix U yields a Hermitian matrix H ; and as this is the case, the KPM can be applied to unitary matrices as well.

3.4 The Choice of Random Unitary Matrices

When generating random unitary matrices for numerical experiments, it is important to ensure that their eigenvalues are distributed uniformly on the unit circle, as predicted by random matrix theory. A common approach is to generate a random

complex matrix A with independent standard normal entries and then construct a unitary matrix from A using either the QR decomposition or the singular value decomposition (SVD). In the QR approach, A is factored as $A = QR$, where Q is unitary and R is upper triangular. The matrix Q is then used as the random unitary matrix. In the SVD approach, A is factored as $A = U\Sigma V^*$, and either U or V (both unitary) can be used as a random unitary matrix. Although both methods produce unitary matrices, their statistical properties differ. The SVD-based approach yields unitary matrices whose eigenvalues are uniformly distributed on the unit circle, matching the theoretical prediction for random unitary matrices (the so-called Haar measure). In contrast, the QR-based approach does not generally produce a uniform distribution of eigenvalues on the unit circle. This distinction is important for applications where the spectral properties of random unitary matrices play a central role, such as in the study of spectral densities. For such purposes, the SVD-based construction is preferable, as it ensures the correct statistical behavior of the eigenvalues.




Figure 3.1: Eigenvalue distributions of random unitary matrices generated via SVD and QR.

Quality Analysis

Let $\tilde{\phi}(t)$ be a regular function, that estimates the spectral density. Since $\phi(t)$ is no function in the classical sense, but all approximations are continuous functions,

$$\left\| \phi(t) - \tilde{\phi}(t) \right\|_{L^p}$$

is not defined, to estimate the error. In the following, we will introduce two methods to circumvent this problem. The first one uses the fact that $\delta(t)$ is a distribution that can be applied to test functions from the Schwartz space $\mathcal{S}(\mathbb{R})$, while the second one regularizes δ -functions by replacing them with continuous and smooth functions. These two approaches are closely related, as we will see later.

4.1 Resolution

In practice, there is often no need to approximate all eigenvalues of A exactly. This would even be highly discontinuous, as we recall the intuitive approach with histograms. The accuracy of the approximation should be oriented towards the desired *resolution*. This means that often only the eigenvalues in a certain subinterval $[a, b] \subset \sigma(A)$ are of interest in a certain context. The size $b - a$ is then called the resolution of the estimate. This insight should be taken as a parameter in our further considerations.

4.2 Restriction on the Schwartz space

In this first method, we consider $\delta(t)$ as a distribution. Let $g \in \mathcal{C}^\infty(\mathbb{R})$ be a test function from the Schwartz space $\mathcal{S}(\mathbb{R})$ defined in Definition 9. Then we have

$$\langle \delta(\cdot - \lambda), g \rangle = \int_{-\infty}^{\infty} \delta(t - \lambda) g(t) dt = g(\lambda)$$

and for all $p, k \in \mathbb{N}_0$

$$\sup_{t \in \mathbb{R}} |t^p g^{(k)}(t)| < \infty$$

The error can therefore be measured as follows:

$$\epsilon_1 = \sup_{g \in \mathcal{S}(\mathbb{R})} \left| \langle \phi, g \rangle - \langle \tilde{\phi}, g \rangle \right|$$

To transfer the concept of resolution to this error, we first consider $\nu_{[a,b]}$ from equation 1.1 and define accordingly

$$\tilde{\nu}_{[a,b]} = \int_a^b n \tilde{\phi}(t) dt$$

with $\tilde{\phi}(t) \in \mathcal{C}^\infty(\mathbb{R})$. Assuming $n = 1$ and thus $\phi(t) = \delta(t)$, an infinite resolution would mean that for arbitrarily small intervals $[a,b]$ the term $|\nu_{[a,b]} - \tilde{\nu}_{[a,b]}|$ also becomes arbitrarily small. Let $a = -\varepsilon, b = \varepsilon$. From the definition of the δ -function it follows that

$$\lim_{\varepsilon \rightarrow 0+} \nu_{[-\varepsilon, \varepsilon]} = 1$$

while for smooth functions $\tilde{\phi}$ it is clear that

$$\lim_{\varepsilon \rightarrow 0+} \tilde{\nu}_{[-\varepsilon, \varepsilon]} = 0$$

This shows that no smooth function converges to the spectral density under continuous increase of the resolution. A carefully balanced approximation would therefore be as precise as a constant one. However, as noted earlier, a finite resolution is often sufficient. We can therefore restrict the Schwartz space $\mathcal{S}(\mathbb{R})$. For example, one could only consider Gaussian distributions of the form

$$g_\sigma(t) = \frac{1}{(2\pi\sigma^2)^{\frac{1}{2}}} e^{-\frac{t^2}{2\sigma^2}}$$

and restrict $\mathcal{S}(\mathbb{R})$ to the subspace

$$\mathcal{S}(\mathbb{R})(\sigma; [\lambda_{us}, \lambda_{os}]) = \{g \mid g(t) \equiv g_\sigma(t - \lambda), \lambda \in [\lambda_{us}, \lambda_{os}]\}$$

Here, λ_{us} and λ_{os} are the upper and lower bounds of the eigenvalues of A as in the previous section, while the parameter σ represents the *target resolution*. We can now use the following metric for quality assessment:

$$E \left[\tilde{\phi}; \mathcal{S}(\mathbb{R})(\sigma; [\lambda_{lb}, \lambda_{ub}]) \right] = \sup_{g \in \mathcal{S}(\mathbb{R})(\sigma; [\lambda_{lb}, \lambda_{ub}])} \left| \langle \phi, g \rangle - \langle \tilde{\phi}, g \rangle \right| \quad (4.1)$$

4.3 Regularizing the Dirac delta distribution

In the second method, we regularize the δ -function. This means that we replace the δ -function with a smooth function $\phi(t)$ that approximates the spectral density. This is done by smoothing the δ -function. For example, we can use the Gaussian distribution with standard deviation σ :

$$\phi_\sigma(t) = \frac{1}{(2\pi\sigma^2)^{\frac{1}{2}}} e^{-\frac{t^2}{2\sigma^2}}$$

This function is smooth and well-defined for all $t \in \mathbb{R}$. It is also normalized, meaning that

$$\int_{-\infty}^{\infty} \phi_\sigma(t) dt = 1$$

This means that $\phi_\sigma(t)$ can be interpreted as a probability density function. The parameter σ controls the width of the Gaussian distribution and thus the resolution of the approximation. This means that the smaller σ is, the more precise the approximation is, but also the more sensitive it is to noise in the data. This is a common approach in numerical analysis to approximate distributions with smooth functions. The function $\phi_\sigma(t)$ is then well-defined and can be used to calculate the error for $p = 1, 2$ and ∞ :

$$\epsilon_2 = \left\| \phi_\sigma(t) - \tilde{\phi}(t) \right\|_p$$

Let

$$\phi_\sigma(t) = \langle \phi(\cdot), g_\sigma(\cdot - t) \rangle = \sum_{j=1}^n g_\sigma(t - \lambda_j)$$

This is the convolution of the spectral density with the Gaussian function g_σ . In similar manner, let

$$\tilde{\phi}_\sigma(t) = \langle \tilde{\phi}(\cdot), g_\sigma(\cdot - t) \rangle$$


Then

$$E \left[\tilde{\phi}; \mathcal{S}(\mathbb{R}) (\sigma; [\lambda_{lb}, \lambda_{ub}]) \right] = \sup_{g \in \mathcal{S}(\mathbb{R}) (\sigma; [\lambda_{lb}, \lambda_{ub}])} \left| \phi_\sigma(t) - \tilde{\phi}_\sigma(t) \right|$$

is the L^∞ error between two well-defined functions. Just like in the other method, σ controls the resolution: the larger the σ , the smoother but also less accurate the approximation. For smaller values, one obtains coarser functions, where small peaks can be detected exactly where the eigenvalues are located. The following graphic is taken from [2, p. 6] and illustrates this effect for four different values of σ : Sei also

$$\phi_\sigma(t) = \langle \phi(\cdot), g_\sigma(\cdot - t) \rangle = \sum_{j=1}^n g_\sigma(t - \lambda_j)$$

Dies ist dann nicht anderes als die "Weichzeichnung" der Spektraldichte durch Gauß-Funktionen der Breite σ .



4.4 The condition of non-negativity

As a probability distribution, the spectral density is non-negative, meaning

$$\forall g \in \mathcal{S}(\mathbb{R}), g \geq 0 : \langle \phi, g \rangle \geq 0$$

Some numerical approximations break this property, including the Kernel Polynomial Method. This leads to large errors and must be considered carefully. Methods to mitigate these negative effects would exceed the scope of this work.

Bibliography

- [1] CHRISTIAN BÄR, *Lineare Algebra und analytische Geometrie*, Springer Spektrum 2018.
- [2] LIN LIN, YOUSSEF SAAD AND CHAO YANG, *Approximating Spectral Densities of Large Matrices*, <https://arxiv.org/pdf/1308.5467v2>
5. Oktober 2014.
- [3] G. H. GOLUB UND C. F. VAN LOAN, *Matrix Computations, 4th Edition*, Johns Hopkins University Press, Baltimore, MD, 3rd ed., 2013.
- [4] R. D. RICHTMYER, *Principles of advanced mathematical physics*, vol 1, Springer Verlag, New York, 1981.
- [5] L. PASTUR, M. SHCHERBINA, *Eigenvalue Distribution of Large Random Matrices*, AMS, 2011.